

# TGS Salt Identification Challenge

## Segment salt deposits beneath the Earth's surface

Several areas of Earth with large accumulations of oil and gas also have huge deposits of salt below the surface.

But unfortunately, knowing where large salt deposits are precisely is very difficult. Professional seismic imaging still requires expert human interpretation of salt bodies. This leads to very subjective, highly variable renderings. More alarmingly, it leads to potentially dangerous situations for oil and gas company drillers.

To create the most accurate seismic images and 3D renderings, TGS (the world's leading geoscience data company) is hoping Kaggle's machine learning community will be able to build an algorithm that automatically and accurately identifies if a subsurface target is salt or not.

**Dataset:** <https://www.kaggle.com/c/tgs-salt-identification-challenge/data>

In [1]:

```
! wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (Windows NT 6.2; WOW64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/80.0.3987.149 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-US,en;q=0.9" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/10151/59042/bundle/archive.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1586444582&Signature=gOljnbjB5OvP4FVSOHsfW5VZ%2F0uYmqQqezFg48aB0eveE64U3zZzzrgAwLp4tVhC9Vf8YJoOywQvBKS9U6IDfuRGZC8uMC%2BusEyHS%2FqJQSwLoJ9MfuqTPFzx%2B%2BxWOonSprVQ7LaFOU0bqWpCsrF%2Fo9Q2EtETvKEfq0aAzzoKBc9N7%2BOg7288FUsJXjJwjZSKwpT3DnACPuUbSqF%2BcuuEMYLvF1WyMH%2FvQZifw2P2YdBYKBXDztDI3TY%2FC4cfOYubKPPaHClabXGX1nxazJhtFkqerDvP2doqSxh6rxuWufYF8JY4JD9eazvaffv3%2FM6ZKgk0aA7tB%2F7mgmsIEbZXzg%3D%3D&response-content-disposition=attachment%3B+filename%3Dtgs-salt-identification-challenge.zip" -c -O 'tgs-salt-identification-challenge.zip'
```

```
--2020-04-07 08:01:02-- https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/10151/59042/bundle/archive.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1586444582&Signature=gOljnbjB5OvP4FVSOHsfW5VZ%2F0uYmqQqezFg48aB0eveE64U3zZzzrgAwLp4tVhC9Vf8YJoOywQvBKS9U6IDfuRGZC8uMC%2BusEyHS%2FqJQSwLoJ9MfuqTPFzx%2B%2BxWOonSprVQ7LaFOU0bqWpCsrF%2Fo9Q2EtETvKEfq0aAzzoKBc9N7%2BOg7288FUsJXjJwjZSKwpT3DnACPuUbSqF%2BcuuEMYLvF1WyMH%2FvQZifw2P2YdBYKBXDztDI3TY%2FC4cfOYubKPPaHClabXGX1nxazJhtFkqerDvP2doqSxh6rxuWufYF8JY4JD9eazvaffv3%2FM6ZKgk0aA7tB%2F7mgmsIEbZXzg%3D%3D&response-content-disposition=attachment%3B+filename%3Dtgs-salt-identification-challenge.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 108.177.97.128, 2404:6800:4008:c04::80
Connecting to storage.googleapis.com (storage.googleapis.com)|108.177.97.128|:443... connected.
HTTP request sent, awaiting response... 416 Requested range not satisfiable
```

The file is already fully retrieved; nothing to do.

In [0]:

```
import zipfile
with zipfile.ZipFile("/content/tgs-salt-identification-challenge.zip", "r") as f:
    f.extractall()
```

In [0]:

```
with zipfile.ZipFile('/content/flamingo.zip', 'r') as f:
    f.extractall('/content/flamingo')
```

In [0]:

```
with zipfile.ZipFile('/content/train.zip') as f:
    f.extractall('/content/train')
```

In [0]:

```
with zipfile.ZipFile('/content/test.zip') as f:
    f.extractall('/content/test')
```

In [0]:

```
import os
import random
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
plt.style.use("ggplot")
```

```

from tqdm import notebook, trange
from itertools import chain
from skimage.io import imread, imshow, concatenate_images
from skimage.transform import resize
from skimage.morphology import label
from sklearn.model_selection import train_test_split

import tensorflow as tf

from tensorflow.keras.models import Model, load_model
from tensorflow.keras.layers import Input, Conv2D, Conv2DTranspose, BatchNormalization, Activation, \
    Dense, Dropout, MaxPooling2D, GlobalMaxPool2D, Lambda, \
    RepeatVector, Reshape, concatenate, add
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau
from tensorflow.keras.optimizers import Adam, RMSprop
from tensorflow.keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img

```

In [7]:

```

# This has three components ('/content/train/images', [], ['64e79513a3.png',...4000 images])

train_images= next(os.walk('/content/train/images'))
print(len(train_images))

# Images are in 3rd position hence train_images[2] / train_masks[2]
print('No of images in train data:', len(train_images[2]))

train_masks= next(os.walk('/content/train/masks'))
print('No of masks in train data:', len(train_masks[2]))

```

3  
No of images in train data: 4000  
No of masks in train data: 4000

In [8]:

```

df= pd.read_csv(r'/content/train.csv')
print(df.shape)
print(df.head())

```

```

(4000, 2)
      id                rle_mask
0  575d24d81d                NaN
1  a266a2a9df           5051 5151
2  75efad62c1  9 93 109 94 210 94 310 95 411 95 511 96 612 96...
3  34e51dba6a  48 54 149 54 251 53 353 52 455 51 557 50 659 4...
4  4875705fb0 1111 1 1212 1 1313 1 1414 1 1514 2 1615 2 1716...

```

In [9]:

```

# We observe zero nulls in 'id' column but 1562 NaNs in 'rle_mask' column
df.isnull().sum()

```

Out[9]:

```

id      0
rle_mask 1562
dtype: int64

```

In [10]:

```

# we find the extension is '.png'
print(train_images[2][0])
print(train_masks[2][0])

```

f409d55fee.png  
f409d55fee.png

Below code can be used to visualize the images and corresponding masks

In [11]:

```

# We now check sample images

```

```

path = "/content/train/images/" # the path is provided

```

```
path= '/content/train/images/' # the path is provided
key= train_images[2][:3]
```

```
path_mask= '/content/train/masks/' # the path is provided
key_mask= train_masks[2][:3]
```

```
image= []
title= []
```

```
for i in key[:3]:
    image.append(path+i)
    title.append(i)
```

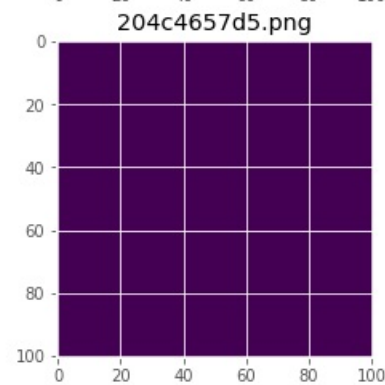

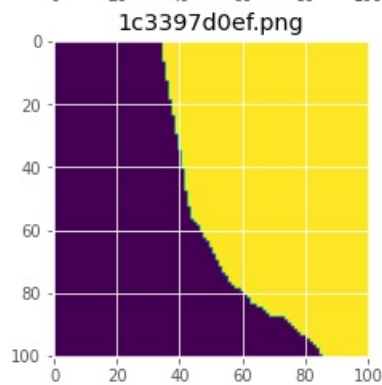
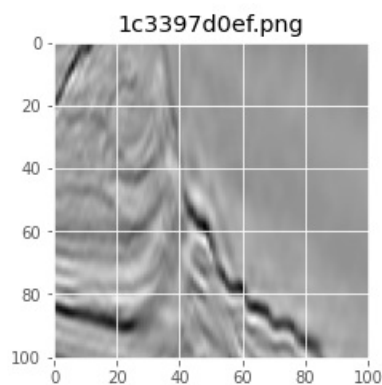
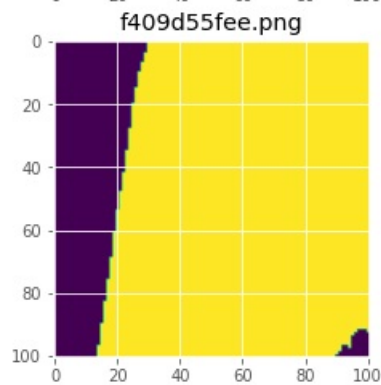
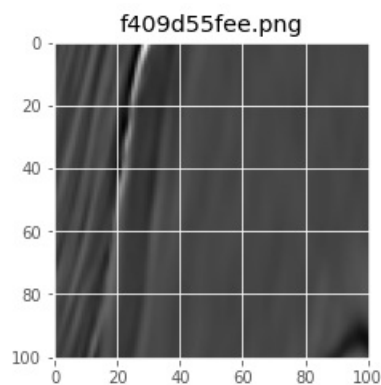
```
for i in key_mask[:3]:
    image.append(path_mask+i)
    title.append(i)
```

```
plt.figure(figsize= (16, 8))
```

```
for i in range(6):
    plt.subplot(2, 3, i+1)
    j= plt.imread(image[i])
    print(j.shape)
    plt.imshow(j)
    plt.title(title[i])
```

```
plt.show()
```

```
(101, 101, 3)
(101, 101, 3)
(101, 101, 3)
(101, 101)
(101, 101)
(101, 101)
```



In [12]:

```
# check the min and max sizes of images
```

```
import time
start= time.time()
```

```
# We check the max and min shape of images
```

```
width=[]
height=[]
for i in train_images[2][:]:
    j= plt.imread(path+i)
    w, h, channel= j.shape
    width.append(w)
    height.append(h)
```

```
i=(M - Width) * (width)
```

```
print('Max Width: ', max(width))
print('Min Width: ', min(width))
print('Max Height: ', max(height))
print('Min Height: ', min(height))
print("Time Taken is: " + str(time.time() - start))
```

Max Width: 101  
Min Width: 101  
Max Height: 101  
Min Height: 101  
Time Taken is: 3.57474946975708

### Load the images and masks into arrays

In [0]:

```
# Reshape the image from 101, 101, 3 to 128, 128, 1
im_height= 128
im_width= 128

X = np.zeros((4000, im_height, im_width, 1), dtype=np.float32)
y = np.zeros((4000, im_height, im_width, 1), dtype=np.float32)
```

In [0]:

```
for i, j in enumerate(train_images[2][:]):
    # Load images
    img = load_img(path+j, color_mode= 'grayscale')
    img = img_to_array(img)
    img = resize(img, (128, 128, 1), mode = 'constant', preserve_range = True)
    # Save images
    X[i] = img / 255.0

for i, j in enumerate(train_masks[2][:]):
    # Load masks
    mask = img_to_array(load_img(path_mask+j, color_mode= 'grayscale'))
    mask = resize(mask, (128, 128, 1), mode = 'constant', preserve_range = True)
    # Save images
    y[i] = mask / 255.0
```

In [16]:

```
print('Shape of X:', X.shape)
print('Shape of Y:', y.shape)
```

Shape of X: (4000, 128, 128, 1)  
Shape of Y: (4000, 128, 128, 1)

In [17]:

```
# Split train and valid
X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size=0.1, random_state=42)

print('Shape of X_train:', X_train.shape)
print('Shape of y_train:', y_train.shape)
```

Shape of X\_train: (3600, 128, 128, 1)  
Shape of y\_train: (3600, 128, 128, 1)

In [18]:

```
print('Shape of X:', X_valid.shape)
print('Shape of Y:', y_valid.shape)
```

Shape of X: (400, 128, 128, 1)  
Shape of Y: (400, 128, 128, 1)

In [0]:

```
# To define a conv2D block (2 times conv2D)

def conv2d_block(input_tensor, filters, kernel_size = 3, batchnorm = True):
```

```

"""Function to add 2 convolutional layers with the parameters passed to it"""
# first layer
x = Conv2D(filters = filters, kernel_size = (kernel_size, kernel_size), activation= 'relu', \
padding = 'same')(input_tensor)
if batchnorm:
    x = BatchNormalization()(x)

# second layer
x = Conv2D(filters = filters, kernel_size = (kernel_size, kernel_size), activation= 'relu', \
padding = 'same')(input_tensor)
if batchnorm:
    x = BatchNormalization()(x)

return x

```

## U-NET Architecture

- Encoder part => (16, 16), (32, 32), (64, 64), (128, 128), 256
- Decoder part => [128, (256, 256), 128], [64, (128, 128), 64], [32, (64, 64), 32], [16, (32, 32), 16], 1

In [0]:

```

def get_unet(input_img, filters = 16, dropout = 0.1, batchnorm = True):
    """Function to define the UNET Model"""
    # Contracting Path
    c1 = conv2d_block(input_img, filters * 1, kernel_size = 3, batchnorm = batchnorm)
    p1 = MaxPooling2D((2, 2))(c1)
    p1 = Dropout(dropout)(p1)

    c2 = conv2d_block(p1, filters * 2, kernel_size = 3, batchnorm = batchnorm)
    p2 = MaxPooling2D((2, 2))(c2)
    p2 = Dropout(dropout)(p2)

    c3 = conv2d_block(p2, filters * 4, kernel_size = 3, batchnorm = batchnorm)
    p3 = MaxPooling2D((2, 2))(c3)
    p3 = Dropout(dropout)(p3)

    c4 = conv2d_block(p3, filters * 8, kernel_size = 3, batchnorm = batchnorm)
    p4 = MaxPooling2D((2, 2))(c4)
    p4 = Dropout(dropout)(p4)

    c5 = conv2d_block(p4, filters * 16, kernel_size = 3, batchnorm = batchnorm)

    # Expansive Path
    u6 = Conv2DTranspose(filters * 8, (3, 3), strides = (2, 2), padding = 'same')(c5)
    u6 = concatenate([u6, c4])
    u6 = Dropout(dropout)(u6)
    c6 = conv2d_block(u6, filters * 8, kernel_size = 3, batchnorm = batchnorm)

    u7 = Conv2DTranspose(filters * 4, (3, 3), strides = (2, 2), padding = 'same')(c6)
    u7 = concatenate([u7, c3])
    u7 = Dropout(dropout)(u7)
    c7 = conv2d_block(u7, filters * 4, kernel_size = 3, batchnorm = batchnorm)

    u8 = Conv2DTranspose(filters * 2, (3, 3), strides = (2, 2), padding = 'same')(c7)
    u8 = concatenate([u8, c2])
    u8 = Dropout(dropout)(u8)
    c8 = conv2d_block(u8, filters * 2, kernel_size = 3, batchnorm = batchnorm)

    u9 = Conv2DTranspose(filters * 1, (3, 3), strides = (2, 2), padding = 'same')(c8)
    u9 = concatenate([u9, c1])
    u9 = Dropout(dropout)(u9)
    c9 = conv2d_block(u9, filters * 1, kernel_size = 3, batchnorm = batchnorm)

    outputs = Conv2D(1, (1, 1), activation='sigmoid')(c9)
    model = Model(inputs=[input_img], outputs=[outputs])
    return model

```

In [33]:

```

input_img = Input((im_height, im_width, 1), name='img')
model = get_unet(input_img, filters=16, dropout=0.05, batchnorm=True)
model.compile(optimizer=Adam(), loss="binary_crossentropy", metrics=["accuracy"])
model.summary()

```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
--------------	--------------	---------	--------------

Layer (type)	Output Shape	Param #	Connected to
img (InputLayer)	[(None, 128, 128, 1) 0		
conv2d_20 (Conv2D)	(None, 128, 128, 16) 160		img[0][0]
batch_normalization_19 (Batch Normalization)	(None, 128, 128, 16) 64		conv2d_20[0][0]
max_pooling2d_4 (MaxPooling2D)	(None, 64, 64, 16) 0		batch_normalization_19[0][0]
dropout_8 (Dropout)	(None, 64, 64, 16) 0		max_pooling2d_4[0][0]
conv2d_22 (Conv2D)	(None, 64, 64, 32) 4640		dropout_8[0][0]
batch_normalization_21 (Batch Normalization)	(None, 64, 64, 32) 128		conv2d_22[0][0]
max_pooling2d_5 (MaxPooling2D)	(None, 32, 32, 32) 0		batch_normalization_21[0][0]
dropout_9 (Dropout)	(None, 32, 32, 32) 0		max_pooling2d_5[0][0]
conv2d_24 (Conv2D)	(None, 32, 32, 64) 18496		dropout_9[0][0]
batch_normalization_23 (Batch Normalization)	(None, 32, 32, 64) 256		conv2d_24[0][0]
max_pooling2d_6 (MaxPooling2D)	(None, 16, 16, 64) 0		batch_normalization_23[0][0]
dropout_10 (Dropout)	(None, 16, 16, 64) 0		max_pooling2d_6[0][0]
conv2d_26 (Conv2D)	(None, 16, 16, 128) 73856		dropout_10[0][0]
batch_normalization_25 (Batch Normalization)	(None, 16, 16, 128) 512		conv2d_26[0][0]
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 128) 0		batch_normalization_25[0][0]
dropout_11 (Dropout)	(None, 8, 8, 128) 0		max_pooling2d_7[0][0]
conv2d_28 (Conv2D)	(None, 8, 8, 256) 295168		dropout_11[0][0]
batch_normalization_27 (Batch Normalization)	(None, 8, 8, 256) 1024		conv2d_28[0][0]
conv2d_transpose_4 (Conv2DTranspose)	(None, 16, 16, 128) 295040		batch_normalization_27[0][0]
concatenate_4 (Concatenate)	(None, 16, 16, 256) 0		conv2d_transpose_4[0][0] batch_normalization_25[0][0]
dropout_12 (Dropout)	(None, 16, 16, 256) 0		concatenate_4[0][0]
conv2d_30 (Conv2D)	(None, 16, 16, 128) 295040		dropout_12[0][0]
batch_normalization_29 (Batch Normalization)	(None, 16, 16, 128) 512		conv2d_30[0][0]
conv2d_transpose_5 (Conv2DTranspose)	(None, 32, 32, 64) 73792		batch_normalization_29[0][0]
concatenate_5 (Concatenate)	(None, 32, 32, 128) 0		conv2d_transpose_5[0][0] batch_normalization_23[0][0]
dropout_13 (Dropout)	(None, 32, 32, 128) 0		concatenate_5[0][0]
conv2d_32 (Conv2D)	(None, 32, 32, 64) 73792		dropout_13[0][0]
batch_normalization_31 (Batch Normalization)	(None, 32, 32, 64) 256		conv2d_32[0][0]
conv2d_transpose_6 (Conv2DTranspose)	(None, 64, 64, 32) 18464		batch_normalization_31[0][0]
concatenate_6 (Concatenate)	(None, 64, 64, 64) 0		conv2d_transpose_6[0][0] batch_normalization_21[0][0]
dropout_14 (Dropout)	(None, 64, 64, 64) 0		concatenate_6[0][0]
conv2d_34 (Conv2D)	(None, 64, 64, 32) 18464		dropout_14[0][0]
batch_normalization_33 (Batch Normalization)	(None, 64, 64, 32) 128		conv2d_34[0][0]
conv2d_transpose_7 (Conv2DTranspose)	(None, 128, 128, 16) 4624		batch_normalization_33[0][0]
concatenate_7 (Concatenate)	(None, 128, 128, 32) 0		conv2d_transpose_7[0][0] batch_normalization_19[0][0]
dropout_15 (Dropout)	(None, 128, 128, 32) 0		concatenate_7[0][0]
conv2d_36 (Conv2D)	(None, 128, 128, 16) 4624		dropout_15[0][0]
batch_normalization_35 (Batch Normalization)	(None, 128, 128, 16) 64		conv2d_36[0][0]

conv2d\_37 (Conv2D) (None, 128, 128, 1) 17 batch\_normalization\_35[0][0]

=====

Total params: 1,179,121

Trainable params: 1,177,649

Non-trainable params: 1,472

=====

In [0]:

```
callbacks= [ EarlyStopping(monitor= 'val_loss', patience= 10, verbose= 1),
              ReduceLROnPlateau(monitor= 'val_loss', patience= 5, min_lr= 0.001, verbose= 1),
              ModelCheckpoint(filepath= '/content/model-tgs-salt.h5', monitor= 'val_loss', verbose= 1, save_best_only= True,
                              save_weights_only= True )
            ]
```

In [34]:

```
history= model.fit(x= X_train, y= y_train, batch_size= 32, epochs= 50, verbose= 1, callbacks= callbacks,
                   validation_data= (X_valid, y_valid))
```

Epoch 1/50  
113/113 [=====] - ETA: 0s - loss: 0.4697 - accuracy: 0.8218  
Epoch 00001: val\_loss did not improve from 0.24889  
113/113 [=====] - 5s 46ms/step - loss: 0.4697 - accuracy: 0.8218 - val\_loss: 2.5187 - val\_accuracy: 0.2845 - lr: 0.0010  
Epoch 2/50  
113/113 [=====] - ETA: 0s - loss: 0.3309 - accuracy: 0.8750  
Epoch 00002: val\_loss did not improve from 0.24889  
113/113 [=====] - 5s 41ms/step - loss: 0.3309 - accuracy: 0.8750 - val\_loss: 1.5341 - val\_accuracy: 0.2952 - lr: 0.0010  
Epoch 3/50  
113/113 [=====] - ETA: 0s - loss: 0.2837 - accuracy: 0.8876  
Epoch 00003: val\_loss did not improve from 0.24889  
113/113 [=====] - 5s 41ms/step - loss: 0.2837 - accuracy: 0.8876 - val\_loss: 0.9180 - val\_accuracy: 0.5301 - lr: 0.0010  
Epoch 4/50  
113/113 [=====] - ETA: 0s - loss: 0.2506 - accuracy: 0.8967  
Epoch 00004: val\_loss did not improve from 0.24889  
113/113 [=====] - 5s 41ms/step - loss: 0.2506 - accuracy: 0.8967 - val\_loss: 1.7537 - val\_accuracy: 0.5471 - lr: 0.0010  
Epoch 5/50  
113/113 [=====] - ETA: 0s - loss: 0.2389 - accuracy: 0.8991  
Epoch 00005: val\_loss did not improve from 0.24889  
113/113 [=====] - 5s 41ms/step - loss: 0.2389 - accuracy: 0.8991 - val\_loss: 0.2762 - val\_accuracy: 0.8803 - lr: 0.0010  
Epoch 6/50  
113/113 [=====] - ETA: 0s - loss: 0.2185 - accuracy: 0.9057  
Epoch 00006: val\_loss improved from 0.24889 to 0.21064, saving model to /content/model-tgs-salt.h5  
113/113 [=====] - 5s 42ms/step - loss: 0.2185 - accuracy: 0.9057 - val\_loss: 0.2106 - val\_accuracy: 0.9105 - lr: 0.0010  
Epoch 7/50  
113/113 [=====] - ETA: 0s - loss: 0.2154 - accuracy: 0.9079  
Epoch 00007: val\_loss did not improve from 0.21064  
113/113 [=====] - 5s 41ms/step - loss: 0.2154 - accuracy: 0.9079 - val\_loss: 0.3783 - val\_accuracy: 0.8383 - lr: 0.0010  
Epoch 8/50  
113/113 [=====] - ETA: 0s - loss: 0.2010 - accuracy: 0.9134  
Epoch 00008: val\_loss did not improve from 0.21064  
113/113 [=====] - 5s 41ms/step - loss: 0.2010 - accuracy: 0.9134 - val\_loss: 0.2190 - val\_accuracy: 0.9085 - lr: 0.0010  
Epoch 9/50  
113/113 [=====] - ETA: 0s - loss: 0.1907 - accuracy: 0.9185  
Epoch 00009: val\_loss did not improve from 0.21064  
113/113 [=====] - 5s 41ms/step - loss: 0.1907 - accuracy: 0.9185 - val\_loss: 0.3868 - val\_accuracy: 0.8302 - lr: 0.0010  
Epoch 10/50  
113/113 [=====] - ETA: 0s - loss: 0.1831 - accuracy: 0.9223  
Epoch 00010: val\_loss did not improve from 0.21064  
113/113 [=====] - 5s 41ms/step - loss: 0.1831 - accuracy: 0.9223 - val\_loss: 0.3012 - val\_accuracy: 0.8931 - lr: 0.0010  
Epoch 11/50  
113/113 [=====] - ETA: 0s - loss: 0.1767 - accuracy: 0.9233  
Epoch 00011: val\_loss improved from 0.21064 to 0.17648, saving model to /content/model-tgs-salt.h5  
113/113 [=====] - 5s 41ms/step - loss: 0.1767 - accuracy: 0.9233 - val\_loss: 0.1765 - val\_accuracy: 0.9286 - lr: 0.0010  
Epoch 12/50  
113/113 [=====] - ETA: 0s - loss: 0.1635 - accuracy: 0.9272  
Epoch 00012: val\_loss did not improve from 0.17648  
113/113 [=====] - 5s 41ms/step - loss: 0.1635 - accuracy: 0.9272 - val\_loss: 0.1963 - val\_accuracy: 0.9227 - lr: 0.0010

```

010
Epoch 13/50
113/113 [=====] - ETA: 0s - loss: 0.1579 - accuracy: 0.9315
Epoch 00013: val_loss did not improve from 0.17648
113/113 [=====] - 5s 41ms/step - loss: 0.1579 - accuracy: 0.9315 - val_loss: 0.2231 - val_accuracy: 0.9120 - lr: 0.0
010
Epoch 14/50
113/113 [=====] - ETA: 0s - loss: 0.1487 - accuracy: 0.9340
Epoch 00014: val_loss did not improve from 0.17648
113/113 [=====] - 5s 41ms/step - loss: 0.1487 - accuracy: 0.9340 - val_loss: 0.2542 - val_accuracy: 0.8917 - lr: 0.0
010
Epoch 15/50
113/113 [=====] - ETA: 0s - loss: 0.1429 - accuracy: 0.9363
Epoch 00015: val_loss did not improve from 0.17648
113/113 [=====] - 5s 41ms/step - loss: 0.1429 - accuracy: 0.9363 - val_loss: 0.2211 - val_accuracy: 0.9008 - lr: 0.0
010
Epoch 16/50
113/113 [=====] - ETA: 0s - loss: 0.1336 - accuracy: 0.9396
Epoch 00016: val_loss improved from 0.17648 to 0.16743, saving model to /content/model-tgs-salt.h5
113/113 [=====] - 5s 41ms/step - loss: 0.1336 - accuracy: 0.9396 - val_loss: 0.1674 - val_accuracy: 0.9270 - lr: 0.0
010
Epoch 17/50
113/113 [=====] - ETA: 0s - loss: 0.1354 - accuracy: 0.9379
Epoch 00017: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.1354 - accuracy: 0.9379 - val_loss: 0.2261 - val_accuracy: 0.8886 - lr: 0.0
010
Epoch 18/50
113/113 [=====] - ETA: 0s - loss: 0.1127 - accuracy: 0.9475
Epoch 00018: val_loss did not improve from 0.16743
113/113 [=====] - 5s 40ms/step - loss: 0.1127 - accuracy: 0.9475 - val_loss: 0.2078 - val_accuracy: 0.9306 - lr: 0.0
010
Epoch 19/50
113/113 [=====] - ETA: 0s - loss: 0.1116 - accuracy: 0.9481
Epoch 00019: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.1116 - accuracy: 0.9481 - val_loss: 0.2283 - val_accuracy: 0.9211 - lr: 0.0
010
Epoch 20/50
113/113 [=====] - ETA: 0s - loss: 0.1129 - accuracy: 0.9485
Epoch 00020: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.1129 - accuracy: 0.9485 - val_loss: 0.2023 - val_accuracy: 0.9258 - lr: 0.0
010
Epoch 21/50
113/113 [=====] - ETA: 0s - loss: 0.0941 - accuracy: 0.9554
Epoch 00021: ReduceLROnPlateau reducing learning rate to 0.001.

Epoch 00021: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.0941 - accuracy: 0.9554 - val_loss: 0.2989 - val_accuracy: 0.8798 - lr: 0.0
010
Epoch 22/50
113/113 [=====] - ETA: 0s - loss: 0.0870 - accuracy: 0.9582
Epoch 00022: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.0870 - accuracy: 0.9582 - val_loss: 0.2233 - val_accuracy: 0.9175 - lr: 0.0
010
Epoch 23/50
113/113 [=====] - ETA: 0s - loss: 0.0803 - accuracy: 0.9604
Epoch 00023: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.0803 - accuracy: 0.9604 - val_loss: 0.2059 - val_accuracy: 0.9257 - lr: 0.0
010
Epoch 24/50
113/113 [=====] - ETA: 0s - loss: 0.0780 - accuracy: 0.9617
Epoch 00024: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.0780 - accuracy: 0.9617 - val_loss: 0.3070 - val_accuracy: 0.8993 - lr: 0.0
010
Epoch 25/50
113/113 [=====] - ETA: 0s - loss: 0.0809 - accuracy: 0.9602
Epoch 00025: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.0809 - accuracy: 0.9602 - val_loss: 0.2360 - val_accuracy: 0.9257 - lr: 0.0
010
Epoch 26/50
113/113 [=====] - ETA: 0s - loss: 0.0731 - accuracy: 0.9638
Epoch 00026: ReduceLROnPlateau reducing learning rate to 0.001.

Epoch 00026: val_loss did not improve from 0.16743
113/113 [=====] - 5s 41ms/step - loss: 0.0731 - accuracy: 0.9638 - val_loss: 0.2228 - val_accuracy: 0.9264 - lr: 0.0
010
Epoch 00026: early stopping

```

In [37]:


```

plt.figure(figsize=(8, 8))
plt.title("Learning curve")
plt.plot(history.history["loss"], label="loss")

```



```
plt.plot(history.history['loss'], label='loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.plot(np.argmax(history.history['val_loss']), np.min(history.history['val_loss']), marker='x', color='r', label='best model')
plt.xlabel("Epochs")
plt.ylabel("log_loss")
plt.legend()
plt.show()
```



## Inference

In [0]:

```
# load the best model
model.load_weights('/content/model-tgs-salt.h5')
```

In [40]:

```
score = model.evaluate(x = X_valid, y = y_valid)
print("Test Loss: ", score[0])
print()
print("Test Accuracy: ", score[1])
```

13/13 [=====] - 0s 12ms/step - loss: 0.1674 - accuracy: 0.9270

Test Loss: 0.16743120551109314

Test Accuracy: 0.9269952178001404

In [41]:

```
train_pred= model.predict(x= X_train, verbose= 1)
val_pred= model.predict(x= X_valid)
```

113/113 [=====] - 1s 13ms/step

In [0]:

```
# Threshold predictions
train_pred_t= (train_pred > 0.5).astype(np.uint8)

val_pred_t= (val_pred > 0.5).astype(np.uint8)
```

In [0]:

```
def plot_sample(X, y, preds, binary_preds, ix=None):
    """Function to plot the results"""
```

*Function to plot the results*  
**if** ix **is** None:  
 ix = random.randint(0, len(X))

has\_mask = y[ix].max() > 0

fig, ax = plt.subplots(1, 4, figsize=(20, 10))  
 ax[0].imshow(X[ix, ..., 0], cmap='seismic')  
**if** has\_mask:  
 ax[0].contour(y[ix].squeeze(), colors='k', levels=[0.5])  
 ax[0].set\_title('Seismic')

ax[1].imshow(y[ix].squeeze())  
 ax[1].set\_title('Salt')


ax[2].imshow(preds[ix].squeeze(), vmin=0, vmax=1)  
**if** has\_mask:  
 ax[2].contour(y[ix].squeeze(), colors='k', levels=[0.5])  
 ax[2].set\_title('Salt Predicted')

ax[3].imshow(binary\_preds[ix].squeeze(), vmin=0, vmax=1)  
**if** has\_mask:  
 ax[3].contour(y[ix].squeeze(), colors='k', levels=[0.5])  
 ax[3].set\_title('Salt Predicted binary');

## Predit on Training set and Val set


In [48]:

*# Check if training data looks all right*  
*# ix is any number between (1 - 3600)*  
 plot\_sample(X\_train, y\_train, train\_pred, train\_pred\_t, ix=14)




In [54]:


plot\_sample(X\_train, y\_train, train\_pred, train\_pred\_t, ix= 114)



In [57]:

plot\_sample(X\_valid, y\_valid, val\_pred, val\_pred\_t, ix= 394)





In [53]:

```
# ix is any number between (1 - 400)
plot_sample(X_valid, y_valid, val_pred, val_pred_t, ix= 30)
```

